

Assignment7

April 26, 2022

0.1 Assignment 7

0.1.1 a.

```
[16]: import os
import json
from pathlib import Path
import gzip
import hashlib
import shutil
import pandas as pd
import pygeohash
import s3fs

# endpoint_url='https://storage.budsc.midwest-datascience.com'
current_dir = Path(os.getcwd()).absolute()
results_dir = current_dir.joinpath('results')

if results_dir.exists():
    shutil.rmtree(results_dir)
results_dir.mkdir(parents=True, exist_ok=True)

def read_jsonl_data():
    #s3 = s3fs.S3FileSystem(
    #    anon=True,
    #    client_kwargs={
    #        'endpoint_url': endpoint_url
    #    }
    # )
    src_data_path = 'data/processed/openflights/routes.jsonl.gz'
    #with s3.open(src_data_path, 'rb') as f_gz:
    with gzip.open(src_data_path, 'rb') as f:
        records = [json.loads(line) for line in f.readlines()]
    return records

def flatten_record(record):
    flat_record = dict()

    for key, value in record.items():
```

```

        if key in ['airline', 'src_airport', 'dst_airport']:
            if isinstance(value, dict):
                for child_key, child_value in value.items():
                    flat_key = '{}_{}'.format(key, child_key)
                    flat_record[flat_key] = child_value
            else:
                flat_record[key] = value
        return flat_record

def create_flattened_dataset():
    records = read_jsonl_data()
    parquet_path = results_dir.joinpath('routes-flattened.parquet')
    return pd.DataFrame.from_records([flatten_record(record) for record in
    ↪records])

```

```

[17]: df = create_flattened_dataset()
df['key'] = df['src_airport_iata'].astype(str) + df['dst_airport_iata'].
    ↪astype(str) + df['airline_iata'].astype(str)

```

```

[18]: partitions = (
        ('A', 'A'), ('B', 'B'), ('C', 'D'), ('E', 'F'),
        ('G', 'H'), ('I', 'J'), ('K', 'L'), ('M', 'M'),
        ('N', 'N'), ('O', 'P'), ('Q', 'R'), ('S', 'T'),
        ('U', 'U'), ('V', 'V'), ('W', 'X'), ('Y', 'Z')
    )

```

```

[19]: #nan values are causing an issue with key assignment so I am removing them from
    ↪the dataset.
df = df[df['src_airport_iata'].isna() == False]

```

```

[20]: #set kv-key equal to the first letter
df['kv_key'] = df['key'].str[0]

#assign a value from the partitions list of tuples
df['kv_key'] = df['kv_key'].apply(lambda x: [str('-'.join(partition)) for
    ↪partition in partitions if (str(x) >= partition[0]) & (str(x) <=
    ↪partition[1])])

# the result of the previous assignment were lists so here I am converting them
    ↪to strings
df['kv_key'] = [''.join(partition) for partition in df['kv_key']]

#here i'm replacing the partitions that have the same start and end letter with
    ↪a single letter
df['kv_key'] = [partition[0] if partition[0] == partition[2] else partition for
    ↪partition in df['kv_key']]

```

```
[21]: df.to_parquet(
        path='results/kv',
        partition_cols=['kv_key']
    )
```

0.1.2 b.

```
[22]: import hashlib

def hash_key(key):
    m = hashlib.sha256()
    m.update(str(key).encode('utf-8'))
    return m.hexdigest()
```

```
[23]: df['hashed'] = df['key'].apply(lambda x: hash_key(x))
df['hash_key'] = df['hashed'].str[0]
```

```
[24]: df.head()
```

```
[24]:
```

	airline_airline_id	airline_name	airline_alias	airline_iata	\
0	410	Aerocondor	ANA All Nippon Airways	2B	
1	410	Aerocondor	ANA All Nippon Airways	2B	
2	410	Aerocondor	ANA All Nippon Airways	2B	
3	410	Aerocondor	ANA All Nippon Airways	2B	
4	410	Aerocondor	ANA All Nippon Airways	2B	

	airline_icao	airline_callsign	airline_country	airline_active	\
0	ARD	AEROCONDOR	Portugal	True	
1	ARD	AEROCONDOR	Portugal	True	
2	ARD	AEROCONDOR	Portugal	True	
3	ARD	AEROCONDOR	Portugal	True	
4	ARD	AEROCONDOR	Portugal	True	

	src_airport_airport_id	src_airport_name	...	dst_airport_dst	\
0	2965.0	Sochi International Airport	...	N	
1	2966.0	Astrakhan Airport	...	N	
2	2966.0	Astrakhan Airport	...	N	
3	2968.0	Chelyabinsk Balandino Airport	...	N	
4	2968.0	Chelyabinsk Balandino Airport	...	N	

	dst_airport_tz_id	dst_airport_type	dst_airport_source	codeshare	equipment	\
0	Europe/Moscow	airport	OurAirports	False	[CR2]	
1	Europe/Moscow	airport	OurAirports	False	[CR2]	
2	Europe/Moscow	airport	OurAirports	False	[CR2]	
3	Europe/Moscow	airport	OurAirports	False	[CR2]	
4	Asia/Krasnoyarsk	airport	OurAirports	False	[CR2]	

	key	kv_key	hashed \
0	AERKZN2B	A	652cdec02010381f175efe499e070c8cbaac1522bac59a...
1	ASFKZN2B	A	9eea5dd88177f8d835b2bb9cb27fb01268122b635b241a...
2	ASFMRV2B	A	161143856af25bd4475f62c80c19f68936a139f653c1d3...
3	CEKKZN2B	C-D	39aa99e6ae2757341bede9584473906ef1089e30820c90...
4	CEKOV2B	C-D	143b3389bce68eea3a13ac26a9c76c1fa583ec2bd26ea8...

	hash_key
0	6
1	9
2	1
3	3
4	1

[5 rows x 42 columns]

```
[25]: df.to_parquet(
        path='results/hash',
        partition_cols=['hash_key']
    )
```

0.1.3 c.

```
[26]: #get hash for datacenters
datacenters = {}

datacenters['west'] = pygeohash.encode(45.5945645, -121.1786823)
datacenters['central'] = pygeohash.encode(41.1544433, -96.0422378)
datacenters['east'] = pygeohash.encode(39.08344, -77.6497145)

print(datacenters)

{'west': 'c21g6s0rs4c7', 'central': '9z7dnebnj8kb', 'east': 'dqby34cjw922'}
```

```
[27]: def closest_datacenter(latitude, longitude):
        geohash = pygeohash.encode(latitude, longitude)
        dist_dict = {}
        closest_datacenter = ''
        last_distance = None
        for key, value in datacenters.items():
            dist = pygeohash.geohash_approximate_distance(str(geohash), str(value))
            dist_dict[key] = dist
            if (last_distance == None) or (dist < last_distance):
                closest_datacenter = key
                last_distance = dist

        return closest_datacenter
```

```
[28]: df['datacenter'] = df[['src_airport_latitude', 'src_airport_longitude']].
      ↪ apply(lambda x: closest_datacenter(x[0], x[1]), axis=1)
```

```
[29]: df.to_parquet(
      path='results/geo',
      partition_cols=['datacenter']
    )
```

0.1.4 d.

```
[30]: def balance_partitions(keys, num_partitions):
      partitions = []

      #get the ideal number of records in each partition
      partition_size = len(keys) / num_partitions

      #get the count of records for each key
      key_grp_cnts = []
      for key in set(keys):
          occurrences = keys.count(key)
          key_grp_cnts.append(tuple([key, occurrences]))

      key_grp_cnts.sort(key=lambda v: v[0].lower())

      total = 0
      partition_list = []
      #loop through the group counts until you exceed partition_size
      for grp in key_grp_cnts:

          #if the total is 0, then this is the first key in the group
          if total == 0:
              min_grp = grp[0]
              last_group = grp[0]

          #if the incremented total exceeds the ideal partition size, then this_
          ↪ key is the max group and reset the total
          if (total + grp[1]) > partition_size:
              max_grp = last_group
              partition_list.append(tuple([min_grp, max_grp]))
              last_group = grp[0]
              total=0
          else:
              last_group = grp[0]
              total += grp[1]

      #add last partition
      partition_list.append(tuple([min_grp, last_group]))
```

```
return partition_list
```

```
[31]: #start by using a series from the df above as the list of keys  
keys = list(df['airline_name'])  
num_partitions=10
```

```
[32]: print(balance_partitions(keys, num_partitions))
```

```
[('40-Mile Air', 'Air Foyle'), ('Air Greenland', 'Amazzonas'), ('Amerijet  
International', 'China Eastern Airlines'), ('China SSS', 'Eurowings'), ('Excel  
Airways', 'Jet Airways'), ('JetBlue Airways', 'Omni Air International'), ('Onur  
Air', 'Shaheen Air International'), ('Shanghai Airlines', 'TransAsia Airways'),  
('Transavia Holland', 'UTair-Express'), ('Valuair', 'Zoom Airlines')]
```

```
[ ]:
```